Machine Learning Engineer Nanodegree

Capstone Project

Movie Box Office Revenue Prediction with Stacked Gradient Boosting Models

Daniel Tejeda
October 1st, 2019

I. Definition

Project Overview

Given its economic potential, the movie box office revenue prediction is a problem that is being actively researched by data scientist and production houses. The main goal of this project is to develop a machine learning model that predicts the revenue of a movie with public data that would have been available before the movie is released. I will attempt to determine with exploratory data analysis and techniques, which features of the movie are more relevant for its monetary success. A model that proves to be successful on this task could be of high value for the decision-making process of a movie production at different stages or even before the movie gets funding, saving production companies millions of dollars.

The dataset used for this project has been collected from TMDB and published on their [Kaggle competition](). The movie details, credits and keywords have been collected from the TMDB Open API. The dataset contains a range of features such as cast, crew, genre, production company, etc, that will require feature engineering prior modelling.

Problem Statement

Forecasting the financial performance of a movie before its release is generally done using some basic statistical techniques described in [1]. These approaches are very common in practice but they often provide only a coarse estimate before the movie is released. [2].

I selected three popular gradient boosting models as base: XGBoost, CATBoost and LightGBM. I decided on boosting algorithms because they've been proven to be useful in Kaggle competitions with limited training data, training time and little expertise for parameter tuning [3]. I also present a final stacked linear meta-model and its results. This kind of ensemble is also very popular in data science competitions because they can boost predictive accuracy by blending the predictions of multiple models.

Metrics

Root-mean-squared-logarithmic-error (RMSLE) is the main metric for evaluation of all models in this project. This metric is similar to Root-mean-squared-error (RMSE), only calculated in logarithmic scale. To calculate it we use:

$$\sqrt{\frac{1}{n} \sum_{i=1}^{n} (\log(y_i + 1) - \log(\hat{y}_i + 1))^2}$$

The main reason to use this metric instead of RMSE is because of the magnitude of the target variable revenue. Since this value can range up to the millions of dollars and we want to get the relative error without considering magnitude, RMSLE is more suited to this scenario. At some points in the code I use the RMSE function since the target variable and predictions are already converted with the logarithm function, which results in RMSLE.

II. Analysis

Data Exploration

The original dataset contains three files:

- **train.csv**. (3000,23) Training data with 3000 records, 22 features and the target variable
- **test.csv**. (4398,22) Testing data with 4398 records and 22 features.
- **sample_submission.csv**. (4398,2) Sample submission file with 4398 records and one column for the target variable.

The Test dataset does not have the target variable in it because of the nature of Kaggle competitions, so we discarded them and split the Train set 80/10/10 for training, validation (10-fold cross validation) and test.

Input Features

- **ID**. Integer unique id of each movie
- **Belongs_to_collection**. Contains the TMDB Id, Name, Movie Poster and Backdrop URL of a movie in JSON format. You can see the Poster and Backdrop Image like this: https://image.tmdb.org/t/p/original/.
  Example: https://image.tmdb.org/t/p/original//iEhb00TGPucF0b4joM1ieyY026U.jpg
- **Budget**. Budget of a movie in dollars. 0 values mean unknown.
- **Genres**. Contains all the Genres Name & TMDB Id in JSON Format
- **Homepage**. Contains the official homepage URL of a movie.
  Example: http://sonyclassics.com/whiplash/ , this is the homepage of Whiplash movie.
- **Imdb_id**. IMDB id of a movie (string). You can visit the IMDB Page like this: https://www.imdb.com/title/

- **Original_language**. Two digit code of the original language, in which the movie was made. Like: en = English, fr = french.
- **Original_title**. The original title of a movie. Title & Original title may differ, if the original title is not in English.
- **Overview**. Brief description of the movie.
- **Popularity**. Popularity of the movie in float.
- **Poster_path**. Poster path of a movie. You can see the full image like this: https://image.tmdb.org/t/p/original/
- **Production_companies**. All production company name and TMDB id in JSON format of a movie.
- **Production_countries**. Two digit code and full name of the production company in JSON format.
- **Release_date**. Release date of a movie in mm/dd/yy format.
- **Runtime**. Total runtime of a movie in minutes (Integer).
- **Spoken_languages**. Two digit code and full name of the spoken language.
- **Status**. Is the movie released or rumored?
- **Tagline**. Tagline of a movie
- **Title**. English title of a movie
- **Keywords**. TMDB Id and name of all the keywords in JSON format.
- **Cast**. All cast TMDB id, name, character name, gender (1 = Female, 2 = Male) in JSON format
- **Crew**. Name, TMDB id, profile path of various kind of crew members job like Director, Writer, Art, Sound etc.

Target variable

- **Revenue**. Total revenue earned by a movie in dollars.

It was discovered that the original datasets had missing values for budget and revenue. We found those missing values in one of the competition's published kernels and implemented a method to fill them right after we load the data.

After filling the missing values for budget and revenue, we still have missing values to address in data pre-processing.

belongs_to_collection	2396	release_date	0
budget	0	runtime	2
genres	7	spoken_languages	20
homepage	2054	status	0
imdb_id	0	tagline	597
original_language	0	title	0
original_title	0	Keywords	276
overview	8	cast	13
popularity	0	crew	16
poster_path	1	revenue	0
production_companies	156		
production_countries	55		

Out of curiosity, I wanted to know which movies were the top 10 of this dataset:

|	title	revenue	release_date
1126	The Avengers	1519557910	4/25/12
1761	Furious 7	1506249360	4/1/15
2770	Avengers: Age of Ultron	1405403694	4/22/15
684	Beauty and the Beast	1262886337	3/16/17
2322	Transformers: Dark of the Moon	1123746996	6/28/11
906	The Dark Knight Rises	1084939099	7/16/12
2135	Pirates of the Caribbean: On Stranger Tides	1045713802	5/14/11
2562	Finding Dory	1028570889	6/16/16
881	Alice in Wonderland	1025491110	3/3/10
734	Zootopia	1023784195	2/11/16

Exploratory Visualization

My first step was to look at the distribution of the target variable revenue. The following figure (left) shows that the distribution is quite skewed. I used log transformation to deal with this, as shown on the right plot.


Distribution of target variable revenue

My next step was to review the correlation of the three quantitative features and the target variable revenue. This is shown in the following scatter matrix.


Makes sense that budget is the feature that shows more correlation with revenue. Another visualization in the following heatmap.


Correlation between quantitative variables

I also wanted to get a sense of how the dataset was distributed throughout the years. The following figure shows the movies released per year. It clearly shows that from 1921-1976 the amount of data is consistently low and starts to get better in the 90's. The year with the highest number of released movies is 2013 with 141 movies in the dataset.

Some more visualizations of Revenue vs Budget. Mean and total revenue per year and finally the mean runtime by year.
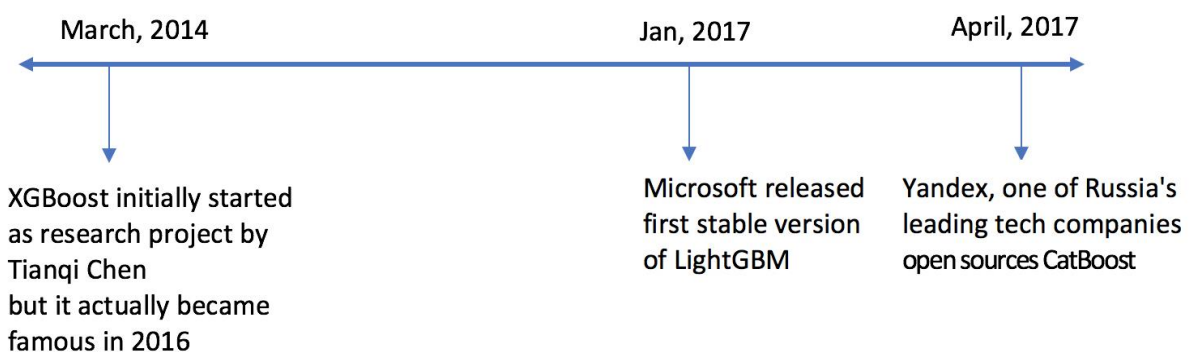
Mean runtime by year

After all these plots were produced, I concluded that runtime seems to be the less correlated variable among the quantiative features to budget.

Algorithms and Techniques

Among the machine learning methods used in practice, Gradient Boosting Decision Tree (GBDT) is one technique that shines in many applications. Tree boosting has been shown to give state-of-the-art results on many standard classification benchmarks [6]

GBDT is an ensemble model of decision trees, which are trained in sequence [5]. In each iteration, GBDT learns the decision trees by fitting the negative gradients (also known as residual errors).

There have been quite a few implementations of GBDT in the literature, here I made use of three very popular ones: XGBoost, LightGBM and CatBoost. The following chronological view shows how recent these algorithms were developed and released.

March, 2014

Jan, 2017

April, 2017

XGBoost initially started as research project by Tianqi Chen but it actually became famous in 2016

Microsoft released first stable version of LightGBM

Yandex, one of Russia's leading tech companies open sources CatBoost

XGBoost, LightGBM and CatBoost release timeline. Source: CatBoost vs. Light GBM vs. XGBoost

Out of the three models, XGBoost is the only one that cannot handle categorical values[[3]]. Another structural difference between them is how they build the decision trees. In many GBDTs, building next tree comprises two steps: choosing the tree structure and setting values in leaves after the tree structure is fixed. To choose the best tree structure, the algorithm enumerates through different splits, builds trees with these splits, sets values in the obtained leaves, scores the trees and selects the best split.[7]. LightGBM uses a novel technique of Gradient-based One-Side Sampling (GOSS) to filter out the data instances for finding a split value while XGBoost uses pre-sorted algorithm & Histogram-based algorithm for computing the best split. In CatBoost the second phase is performed using traditional GBDT scheme and for the first phase uses a modified version.

K-fold validation

I am using 10-fold cross validation for training and validation. This will split the dataset into 10 parts, train on 9 and validate on 1 and repeat for all combinations of train-validation splits.

Stacking

Ensemble methods are commonly used to boost predictive accuracy by combining the predictions of multiple machine learning models. Model stacking is an efficient ensemble method in which the predictions that are generated by using different learning algorithms are used as inputs in a second-level learning algorithm. This second-level algorithm is trained to optimally combine the model predictions to form a final set of predictions [4].

Benchmark

A K-nearest neighbors (KNN) regressor was used as Stage-1 benchmark. Stage-2 benchmark was chosen from the "vanilla" three boosting algorithms, the one with better results: LightGBM.

III. Methodology

Data Preprocessing

Categorical features have a discrete set of values which are not necessarily comparable with each other, making them not suitable to be used in binary decision trees directly. The most widely used technique to approach these, which is usually applied to low-cardinality categorical features is one-hot encoding: the original feature is removed and a new binary variable is added for each category 7. One-hot encoding was done during the preprocessing phase of the data for all models including the benchmark. Another scenario that I found was List/JSON format for some columns. For these, I first extracted the information from the list of JSON objects and then reformatted them.

I created a general method for data preprocessing and divided it into steps methods as follows:

```
def preprocess_data(in_data):

    #copy
    data = in_data.set_index('id')

    data_prep_clean_na(data) # 1. Clean NA
    data_prep_dates(data) # 2. Release Date decomposition
    data_prep_collection(data) # 3. Collection
    data_prep_homepage(data) # 4. Homepage
    data_prep_genres(data) # 5. Genre
    data_prep_cast_crew(data) # 6. Cast and Crew
    data_prep_lang(data) # 7. Original Language
    data_prep_prod_companies_countries(data) # 8. Production Countries and Companies
    data_prep_final(data) # 9. Final cleanup and logarithmic conversion of quantities

    return data
```

1. **Clean NA.** The first step on the data cleansing was to get rid of the missing values. The only quantitative variable with missing values was runtime so I replaced the NA with the mean (only two missing values). I filled all other fields with a blank string.

2. **Release Date decomposition.** The date is in string format with mm/dd/yy. I broke it down into five different features: Day, Month, Year, Dow (day of week), and Quarter. I also completed the year field to be the full 4-digit year.

3. **Collection.** This field contained the information about the collection in JSON format. I decomposed it into one boolean field (0,1) to indicate if the movie belongs to a collection and did one-hot encoding for the ones that belong to a collection with 3 or more movies. Total columns added: 34.

4. **Homepage.** I created a boolean field to indicate if the movie has homepage or not, 946 out of 3000 have homepage.

5. **Genre.** Since one movie can belong to more than one gender, I first extracted the JSON information and put this into another column genre_new with the genre names separated by a pipe. I then did one-hot encoding of the values, adding 20 columns to the dataset.

6. **Cast and Crew.** These two field were decomposed into eight fields based on the gender of people and total count of cast and crew members, i.e. cast_g1 is the total count of all female members of the cast. 1=female, 2=male and 0 is not specified.

7. **Original Language.** Simple full one-hot encoding, 36 columns added.

8. **Production Countries and Companies.** These two columns had the same cardinality as Genre, meaning one movie could be produced by more than one company in more than one country. I applied the same approach of extracting the JSON names for both, put them in a separate column as strings separated by pipes and then one-hot encoded them. This process added too many columns at first, so I put a minimum count of 15 for companies and 5 for countries. The result was 48 columns added for companies and 39 for countries.

9. **Final cleanup and logarithmic conversion of quantities.** In this final method I convert the quantitative input features (budget, runtime and popularity) with logarithm and get rid of the original and temporary features created during data preprocessing.

Implementation

The model implementation is in class MovieRevenuePredictor. This class contains the benchmark, all three GBDT models and the meta-model along with data and methods for training, validating and testing.

Initialization

The class constructor has the following parameters:

- **data**. full dataset
- **random_seed=1.** random seed to be used across all methods to get reproducible results
- **splits=10.** number of training/validation data folds, default 10
- **test_size=0.1.** test size, default 10%
- **gridsearch=False.** added to get the best models through Grid Search. Default is False which means 'vanilla' models are used.

As mentioned, the original Test dataset is not usable because it has no labels. I do the splitting of the dataset into Train-Validation and Test in the method prepare_data of the class.

```python
def prepare_data(self, dataset, splits, test_size):

    train, test,  = train_test_split(dataset, test_size=test_size, random_state=self.random_seed)

    self.data = {
        'raw': dataset,
        'train': train,
        'test' : test
    }

    kfold = KFold(splits, shuffle = True, random_state = self.random_seed)

    self.fold = list(kfold.split(self.data['train'].values))
```

I also created the test-validation splits here and save it into the class attribute fold to be used later on training and meta-feature generation.

After data is ready, we create the base 'vanilla' models. If the parameter gridsearch is set to True, we do a Grid Search on the base models (XGBoost, LightGBM and CatBoost) to find the best parameters.

Training, testing and predicting with base models

You can train the base models (and the benchmark) with the train method. You have a number of parameters to pass, the more relevant ones are models for passing a list of model names (default is all the base models) and stacking which controls whether stacking needs to be prepared. If set to

True, meta-features are generated for each of the validation folds when training the base models and the meta-model is also trained at the end.

For testing and prediction I created methods:

- **Test** can take an additional dataset or do testing on the default one if nothing is specified.
- **Predict** uses the base models to make predictions.

Stacking implementation

I used the train-validation splits to produce the meta-features in a new dataset called train_meta in the following way:

1. For each base model I created a new meta-feature (column) in the train_meta dataset,
2. I made predictions on the test fold only, with the base model trained on the other 9 folds
3. I continued to fill the missing values for each of the validation folds until the meta-feature was complete.

It's important to note that the meta-features in row i of train_meta are not dependent on the target value in row i because they were produced using information that excluded the target_i in the base models' fitting procedure.

Similarly to the base models, I created the methods train_meta, test_meta and predict_meta with the same idea behind them.

Refinement

I originally started using the default hyperparameters for the base models. They provided very good initial scores but after I had the overall structure I created the find_best_params method to have a grid search and results improved. This was a resource intensive task that used almost all 40 cpus in my server.
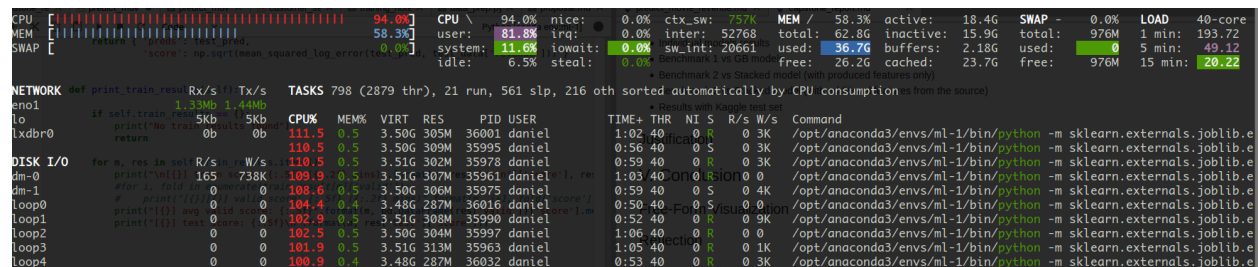
The resulting best models:

```
{'xgb': XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
         colsample_bynode=1, colsample_bytree=1, eval_metric='rmse', gamma=0,
         importance_type='gain', learning_rate=0.05, max_delta_step=0,
         max_depth=30, min_child_weight=6, missing=None, n_estimators=200,
         n_jobs=1, nthread=None, objective='reg:linear', random_state=0,
         reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=681,
         silent=True, subsample=1, verbosity=1),
 'lgb': LGBMRegressor(bagging_seed=681, boosting_type='gbdt', class_weight=None,
         colsample_bytree=1.0, importance_type='split', learning_rate=0.05,
         max_depth=50, metric='rmse', min_child_samples=20,
         min_child_weight=0.001, min_split_gain=0.0, n_estimators=200,
         n_jobs=-1, num_leaves=300, objective='regression',
         random_state=None, reg_alpha=0.0, reg_lambda=0.0, silent=True,
         subsample=1.0, subsample_for_bin=200000, subsample_freq=0,
         use_best_model=True),
 'cat': <catboost.core.CatBoostRegressor at 0x7f6a188a0908>}

movie_pred.best_models['cat'].get_params()

{'eval_metric': 'RMSE',
 'verbose': False,
 'random_seed': 681,
 'loss_function': 'RMSE',
 'depth': 4,
 'iterations': 300,
 'l2_leaf_reg': 4,
 'learning_rate': 0.1}
```

The initial stacking approach (and a very popular one in Kaggle) was to apply fixed weights to the results of the base models. I searched for a better approach and found the one [9] I finally use that involves the a linear regression meta-model with the meta-features produced by the base models.

IV. Results

Model Evaluation and Validation

Since GBDT models have proven their effectiveness already, one of the main points I wanted to evaluate was the different combinations of stacking. I created two different MovieRevenuePredictor objects: one with vanilla models (no hyperparameter tuning) and another one with the tuned parameters as result of the grid search. For each of these I evaluate:

- Benchmark (KNN)

- Individual GBDT models results (XGB, LGB, CAT)
- Stacked combinations of GBDT models (XGB+LGB, XGB+CAT, LGB+CAT)
- Stacked GBDT models (XGB+LGB+CAT)
- Stacking all models (KNN+XGB+LGB+CAT)

I also submitted the results of both the vanilla and tuned models to Kaggle to have another set of scores to compare. To my surprise, even including the KNN in the final meta-model resulted in slightly better test results than just the gradient boosting models alone in both my testing set and in Kaggle's private dataset.

		[knn]	[xgb]	[lgb]	[cat]	[xgb_lgb]	[xgb_cat]	[lgb_cat]	[xgb_lgb_cat]	[knn_xgb_lgb_cat]
Test Dataset	vanilla	2.6557	2.0858	2.01156	2.02935	2.014607	2.049722	1.991545	2.006726	2.006338
	tuned	2.6557	1.93643	1.99794	1.90409	1.926032	1.891558	1.900557	1.891551	1.891119
Kaggle Score	vanilla	2.63532	2.06713	1.96241	2.02272	1.97772	2.02940	1.95684	1.97188	1.97171
	tuned	2.63532	1.88709	1.98514	1.85581	1.86709	1.82993	1.84856	1.82988	1.82961

The Kaggle competition is closed so getting into the public leaderboard is no longer possible. A score of **1.82988** would get me around position 325 according with their final results.

Justification

The final model score was **1.891119**. The benchmark-1 model achieved RMSLE of **2.6557**, so this is a marked improvement. Additionally, all three GBDT models got better scores than the benchmark: 1.93643, 1.99794, 1.90409 in the tuned version of the predictor.

The benchmark-2 (vanilla LightGBM) scored **2.01156** so it's also improved by the final model score. It's worth noting that hyperparameter tuning made CatBoost the best of the three base models with a score of **1.90409** in the test set and **1.85581** in Kaggle.

The final model score as provided by Kaggle was **1.82988** with a different, private dataset. The similarities of the results between my test dataset and Kaggle gave me additional validation of my scores. Overall, the results were satisfactory and solved the problem to a good degree considering we only had 3000 samples of data.

V. Conclusion

Free-Form Visualization

The top 30, most important features for each of the GBDT models are shown in the following charts.

Features Importance [xgb]

Features Importance [lgb]

Features Importance [cat]
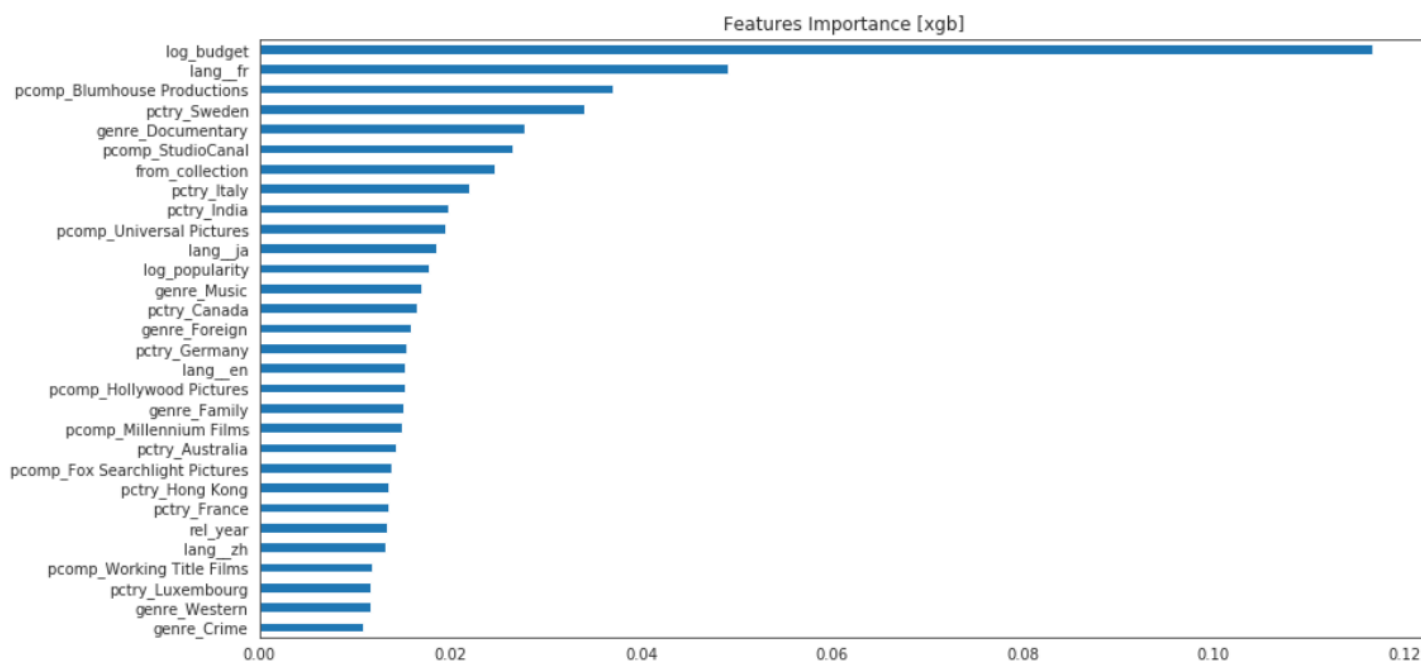
There are evident ones like budget and popularity but I find interesting to see how the algorithms had features like pctry_Sweden in the top features.

Reflection

In conclusion the results using feature engineering and the straightforward application of well-known gradient boosting algorithms gave very good results, even for the default hyperparameters. It was particularly satisfying because I was not sure initially if I could improve the base models with the stacked one.

The initial challenge was to come up with how to manipulate the features and pre-process the data but once I have that and the basic skeleton of the main class, I could start experiment a lot and get good insights. My initial tests always favored XGBoost, but when I reduced the number of one-hot encoded features I started to get better scores with LightGB and CatBoost. Another unexpected result was the addition of KNN to the stack, which I did by mistake and, even though it was minimal, ended up with a better score.

Regarding stacking, the simplicity of the linear regression model made me think a lot about a more complex stacking approach like feeding the base algorithms in a recurrent fashion, but I am not really sure if this would be effective given the fact that the meta features are produced with the base models that are already recurrent.

Improvement

The list of things to improve in machine learning problems is endless. One evident one for this particular case is to get external, public data from other sources, and more data samples since 3000 is a very low amount. I would also try recurrent neural networks for text analysis on the free-text fields that were ignored.

The hyperparameter tuning is also another area of improvement since I didn't try too many combinations given the resources needed. Additionally, LightGBM and CatBoost can both take categorical features and this is listed as one of their strengths but I did not use this option since I wanted to have all three with the same data and XGBoost was not able to handle categorical values.

References

[1] Simonoff, J. S. and Sparrow, I. R. Predicting movie grosses: Winners and losers, blockbusters and sleepers. In Chance, 2000

[2] Matt Vitelli. Predicting Box Office Revenue for Movies

[3] Alvira Swalin, CatBoost vs. Light GBM vs. XGBoost

[4] Sill, J., Takacs, G., Mackey, L., and Lin, D. 2009. Feature-Weighted Linear Stacking

[5] Jerome H Friedman. Greedy function approximation: a gradient boosting machine. 2001

[6] Tianqi Chen and Carlos Guestrin. XGBoost: A Scalable Tree Boosting System

[7] D. Micci-Barreca. A preprocessing scheme for high-cardinality categorical attributes in classification and prediction problems. 2001

[8] Anna Veronika Dorogush, Vasily Ershov and Andrey Gulin. CatBoost: gradient boosting with categorical featuressupport

[9] Ben Gorman. A Kaggler's Guide to Model Stacking in Practice 2016